

System Scalability

- Scalability
 - Vertical & Horizontal Scaling
- Horizontal Scalability
 - Replication
 - Services
 - Caching
 - Asynchronous Process
 - Partitioning
- Load Balancing
 - Load Balancers
 - Service Discovery
 - DNS & Geo Load Balancing
- Micro-Services
 - Architecture
 - Transactions
 - SAGA Pattern
 - NoSQL



Performance Vs Scalability

FIXED LOAD

- Performance
 - Low Latency
 - High Throughput
 - Concurrency
 - Single Machine – Multi-Threading
 - Multi Machine – Multi Threading + Multi-Processing = Distributed Processing
 - Capacity


VARIABLE LOAD

- Scalability
 - High Throughput
 - Ability of a system to increase its throughput by adding more hardware capacity
 - Both ways – UP and DOWN

Vertical & Horizontal Scalability


Vertical

- Easier to achieve
- Limited scalability



Horizontal

- Hard to achieve
- Unlimited scalability



Reverse Proxy

- Client needs to know only about the address of the Reverse Proxy
- Reverse Proxy can also act as a load balancer




System For Scalability Discussion




Scalability Principle

- Decentralization – *Monolith is an anti-pattern for Scalability*
 - More workers – Instances, Threads
 - Specialized workers – Services
- Independence
 - Multiple workers are as good as a single worker if they can't work independently
 - They must work concurrently to maximum extent
 - Independence is impeded by
 - Shared resources
 - Shared mutable data



Modularity

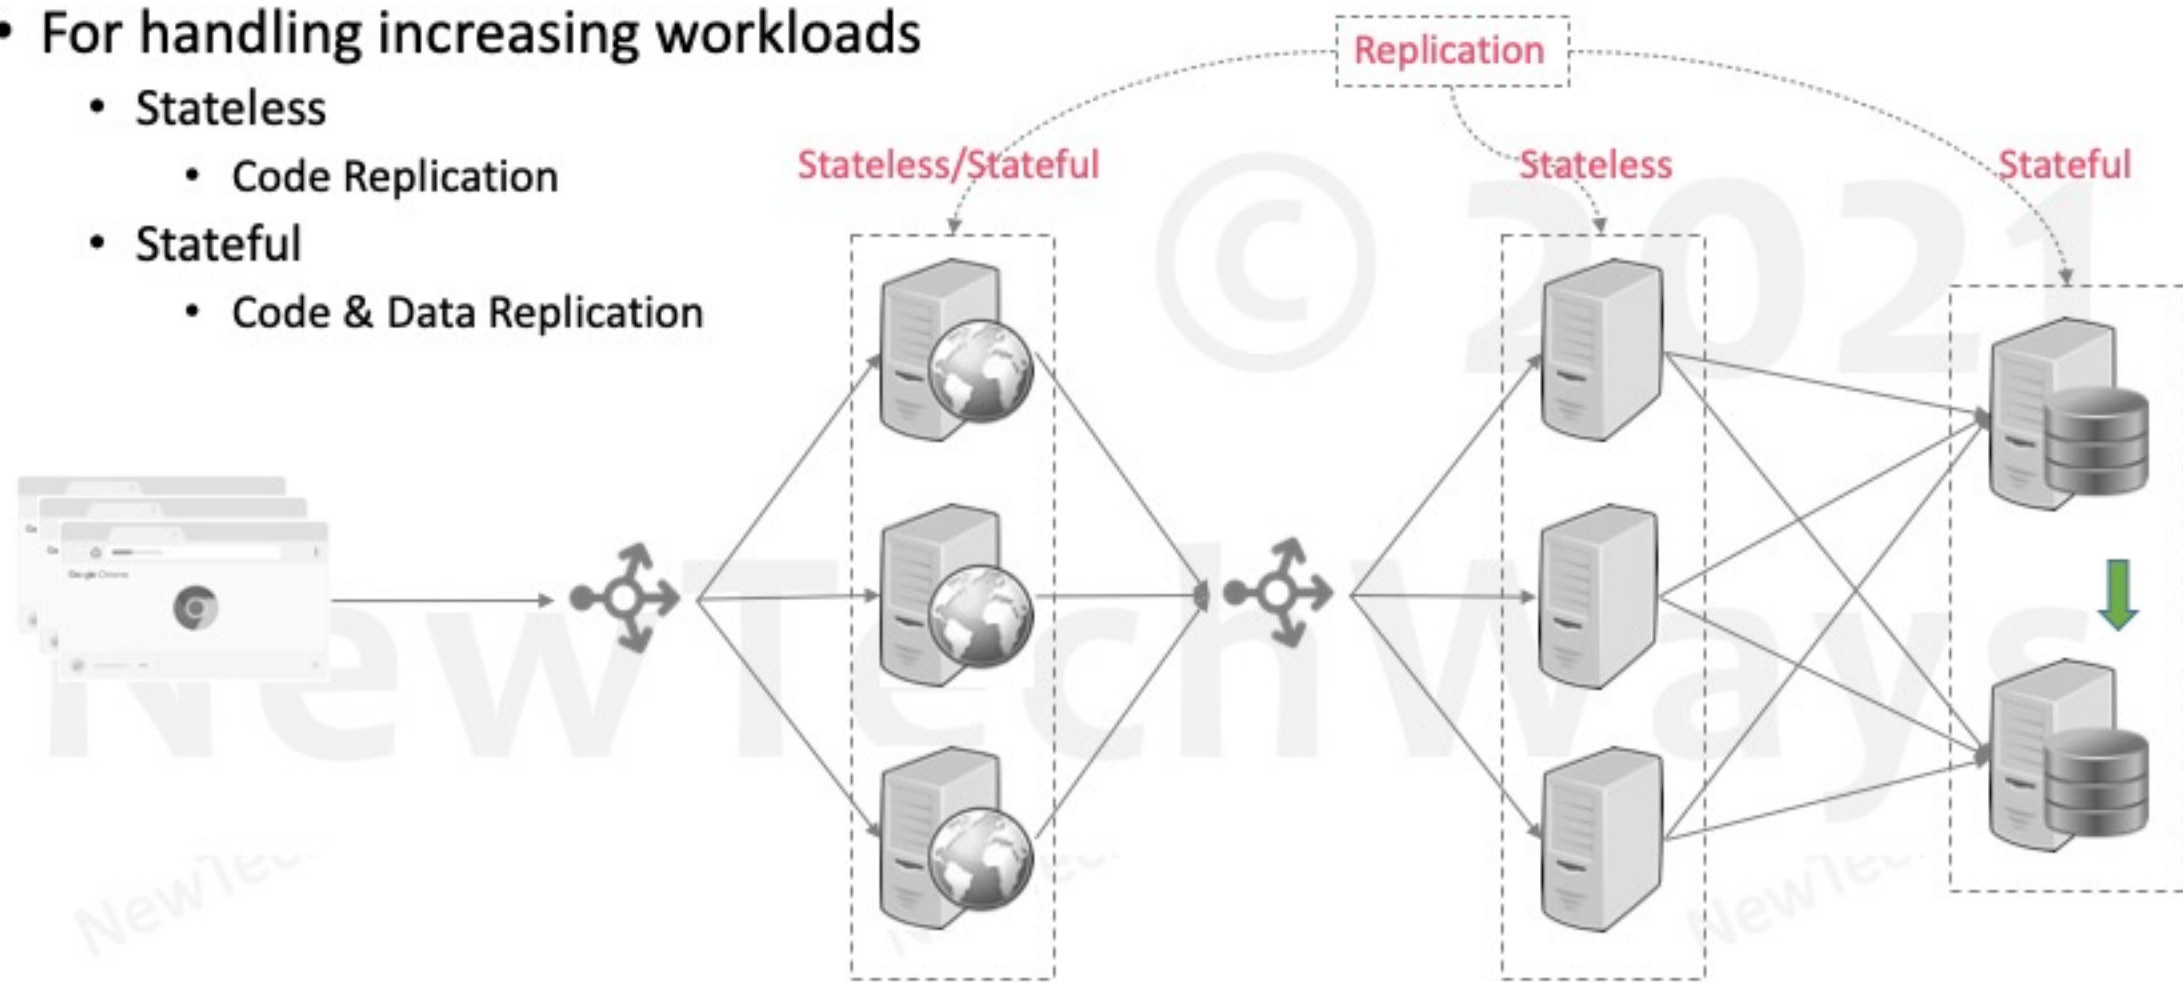
- Scalable architecture starts with modularity
 - Provides the foundation for breaking an application into more specialized functions/services



Replication

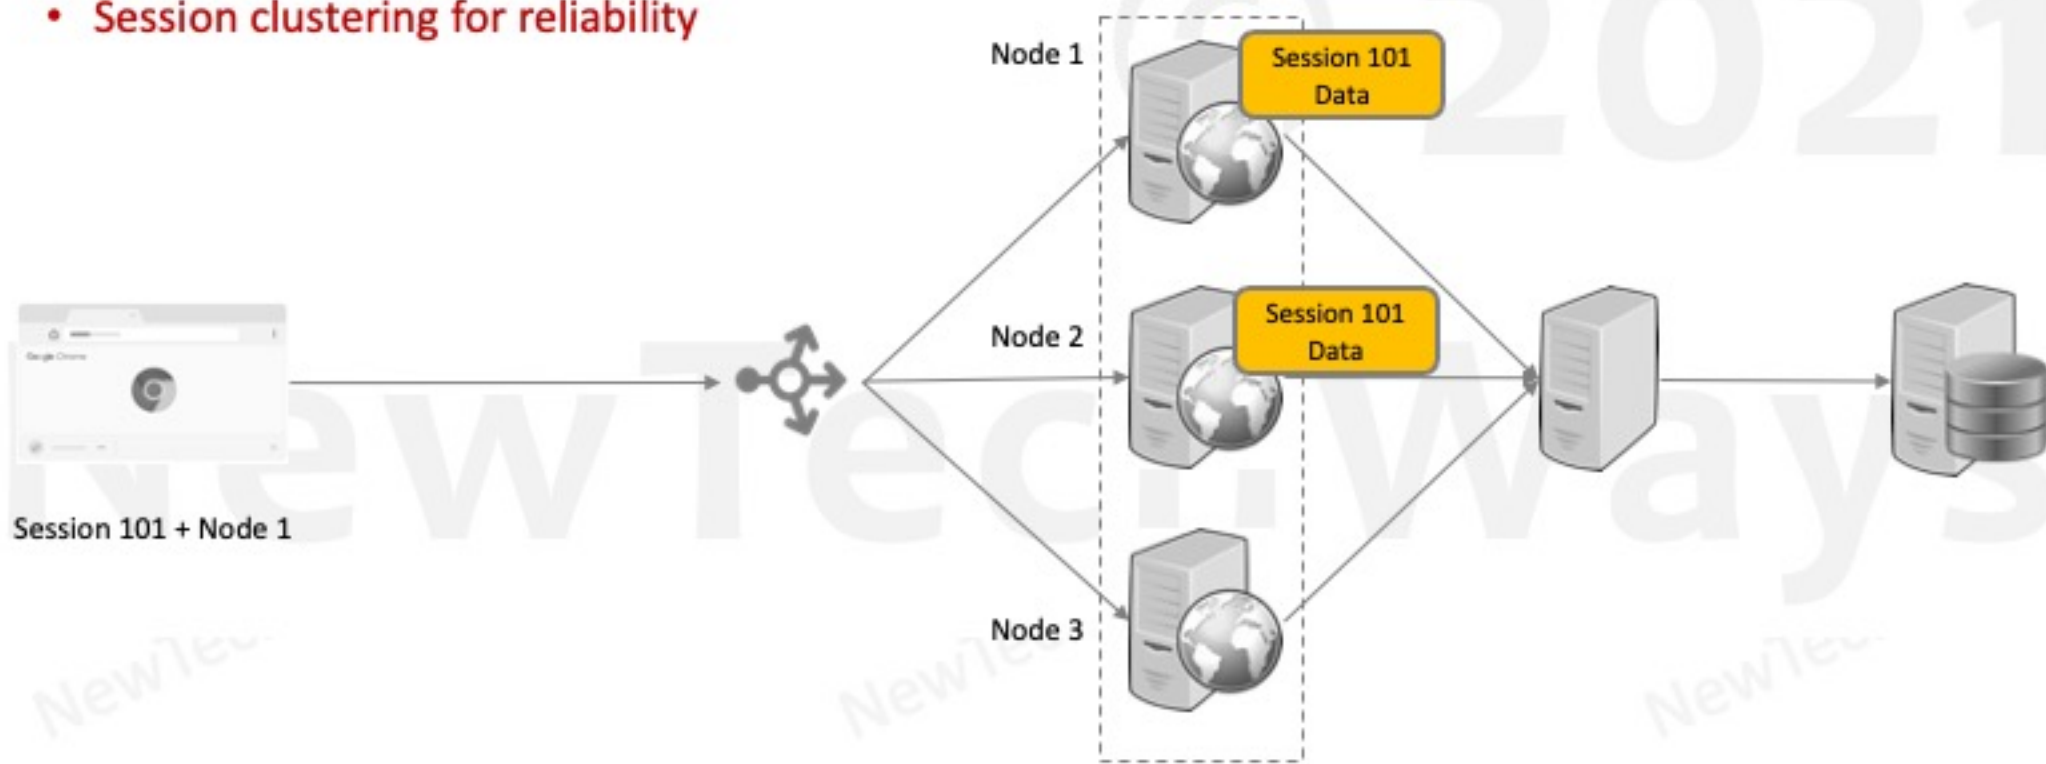
- For handling increasing workloads

- Stateless
 - Code Replication
- Stateful
 - Code & Data Replication



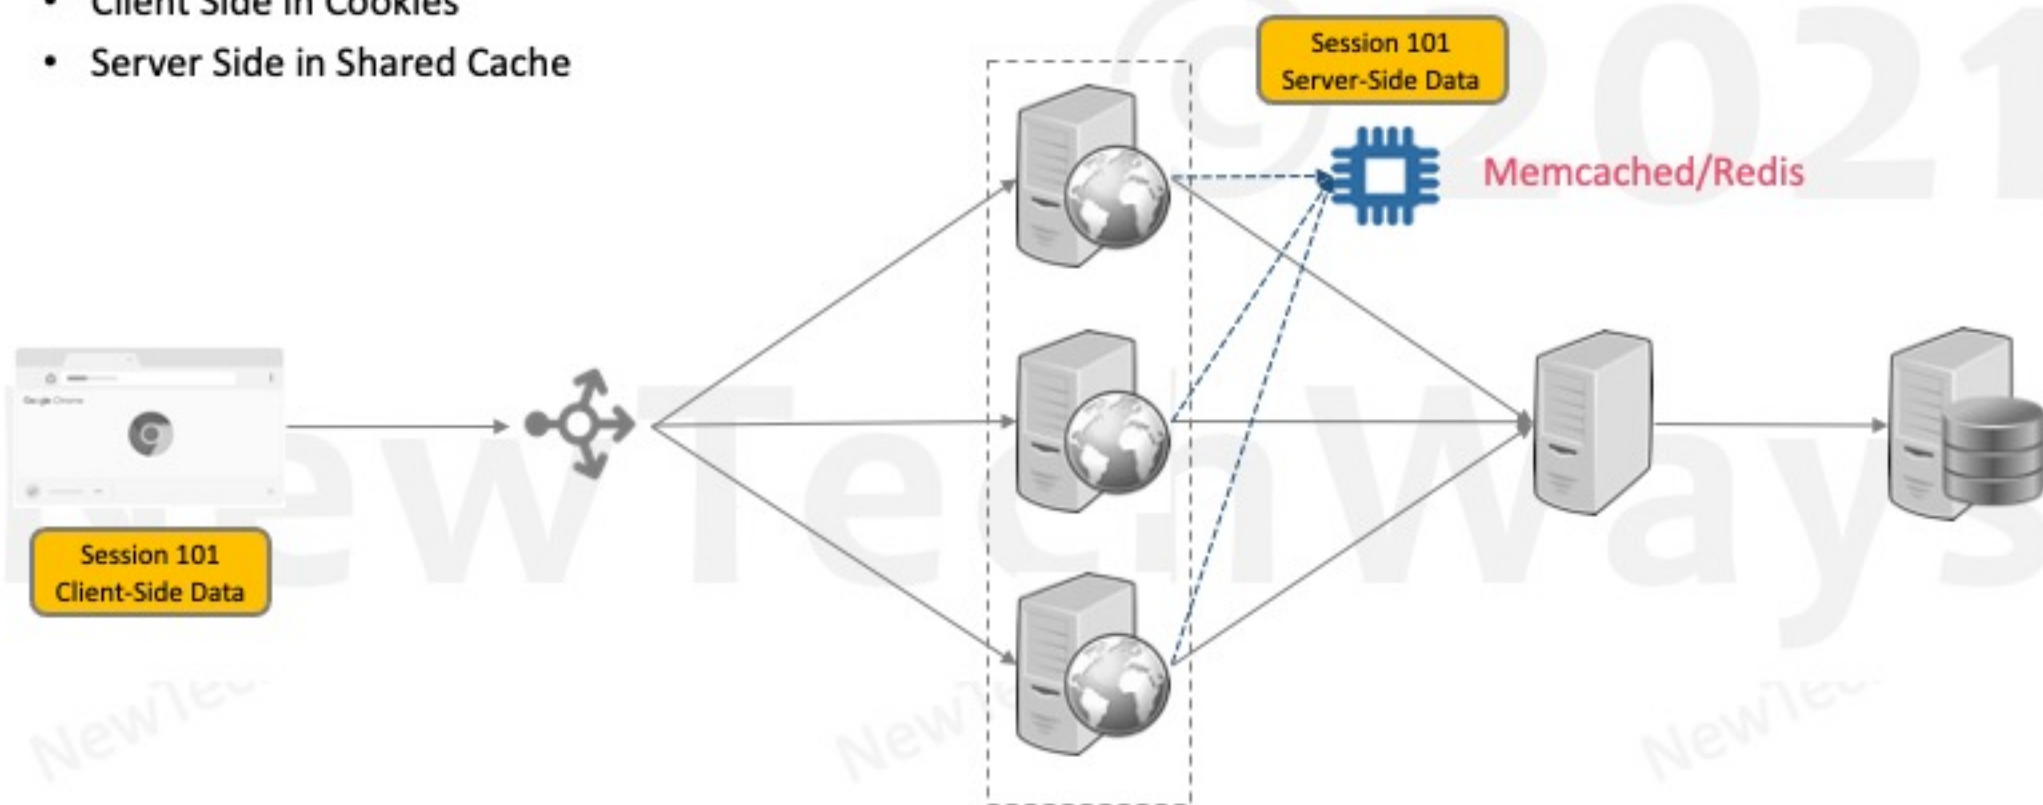
Web Stateful Replication

- When low latency is required
 - Sticky sessions/Session affinity
 - Sessions occupy memory
 - Session clustering for reliability




Web Stateless Replication

- For higher scalability at the expense of higher latency
- Session data can be stored on
 - Client Side in Cookies
 - Server Side in Shared Cache




Service Replication

- Stateless replication – same as web stateless



Database Replication

- For higher read scalability
- For high availability



Database Replication

- Master-Slave (Primary-Secondary)

- Asynchronous

- Low latency writes
 - Eventually Consistent
 - Data Loss

- Synchronous

- Consistent
 - High latency writes
 - Low write availability


- High Read Scalability
 - High Read Availability
 - No Write Conflicts

- Master-Master (No-Master/Peer-To-Peer)

- Asynchronous

- Write conflicts
 - High availability

- High Read Scalability
 - High Read Write Availability
 - Transaction ordering issues



Need For Specialized Services

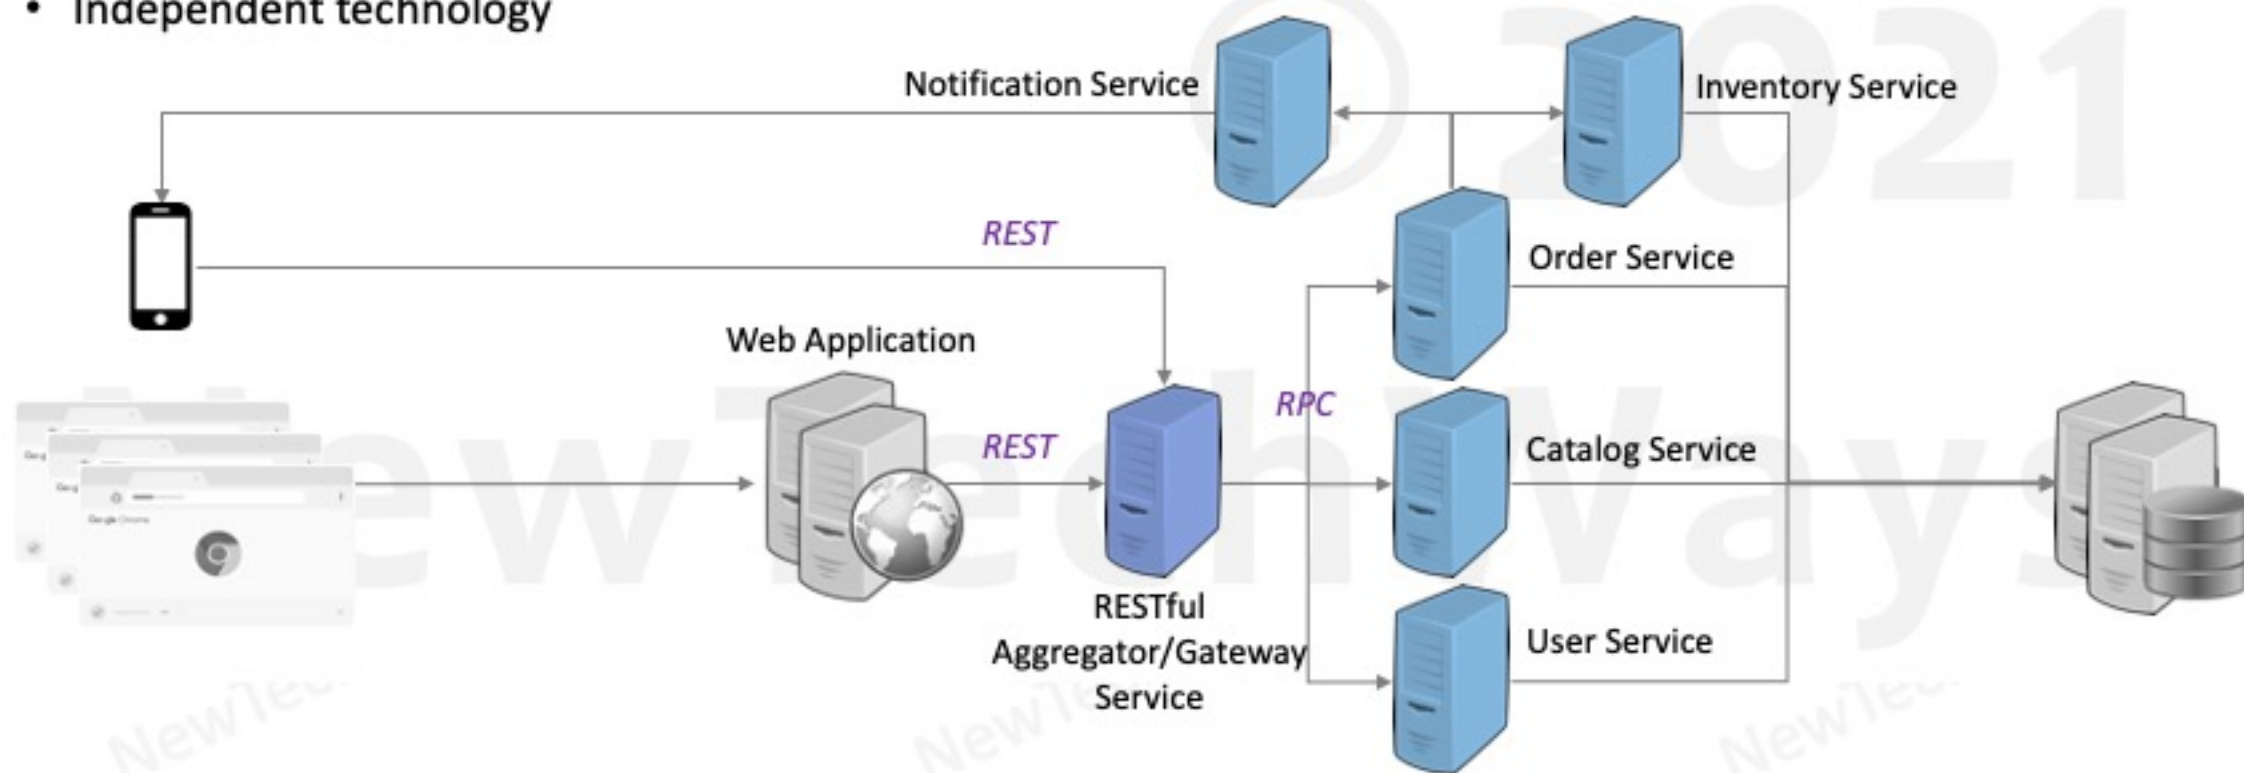
Service Modules:

- User
- Catalog
- Order
- Inventory
- Notification



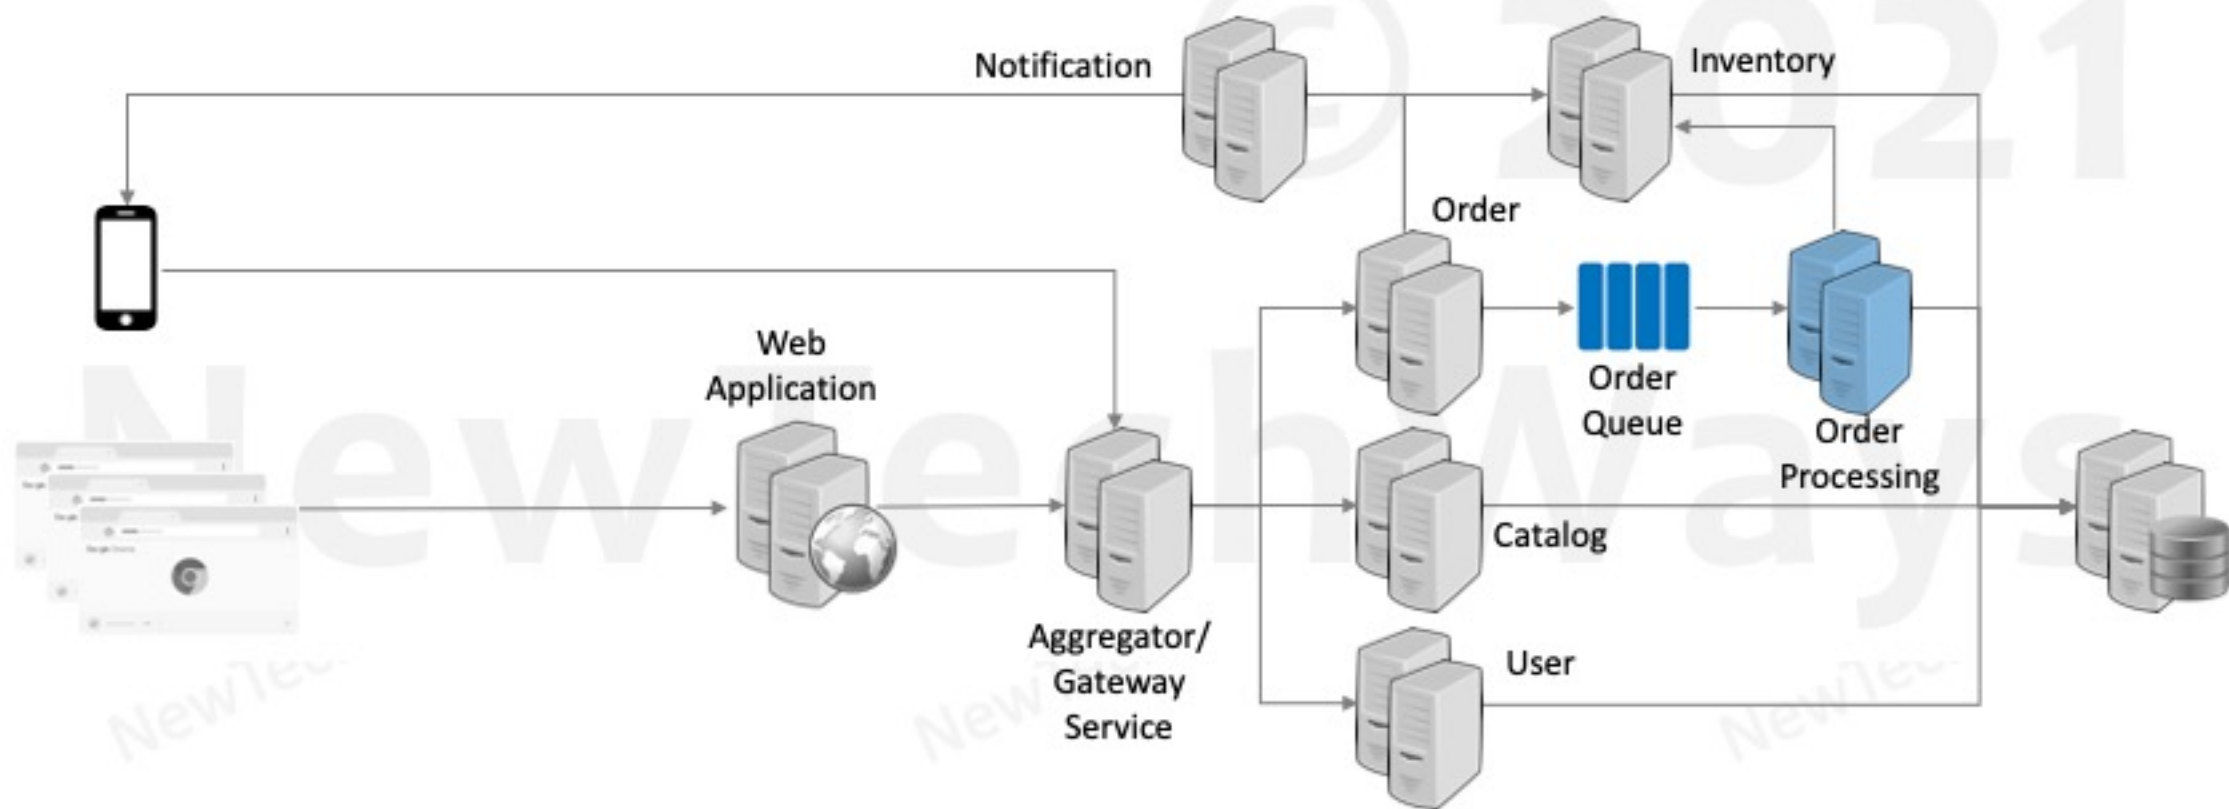
Specialized Services (SOAP/REST-Services)

- Partially independent development and deployment
- Independent scalability
- Independent technology




Asynchronous Services

- Async services effectively reduces write load from a database



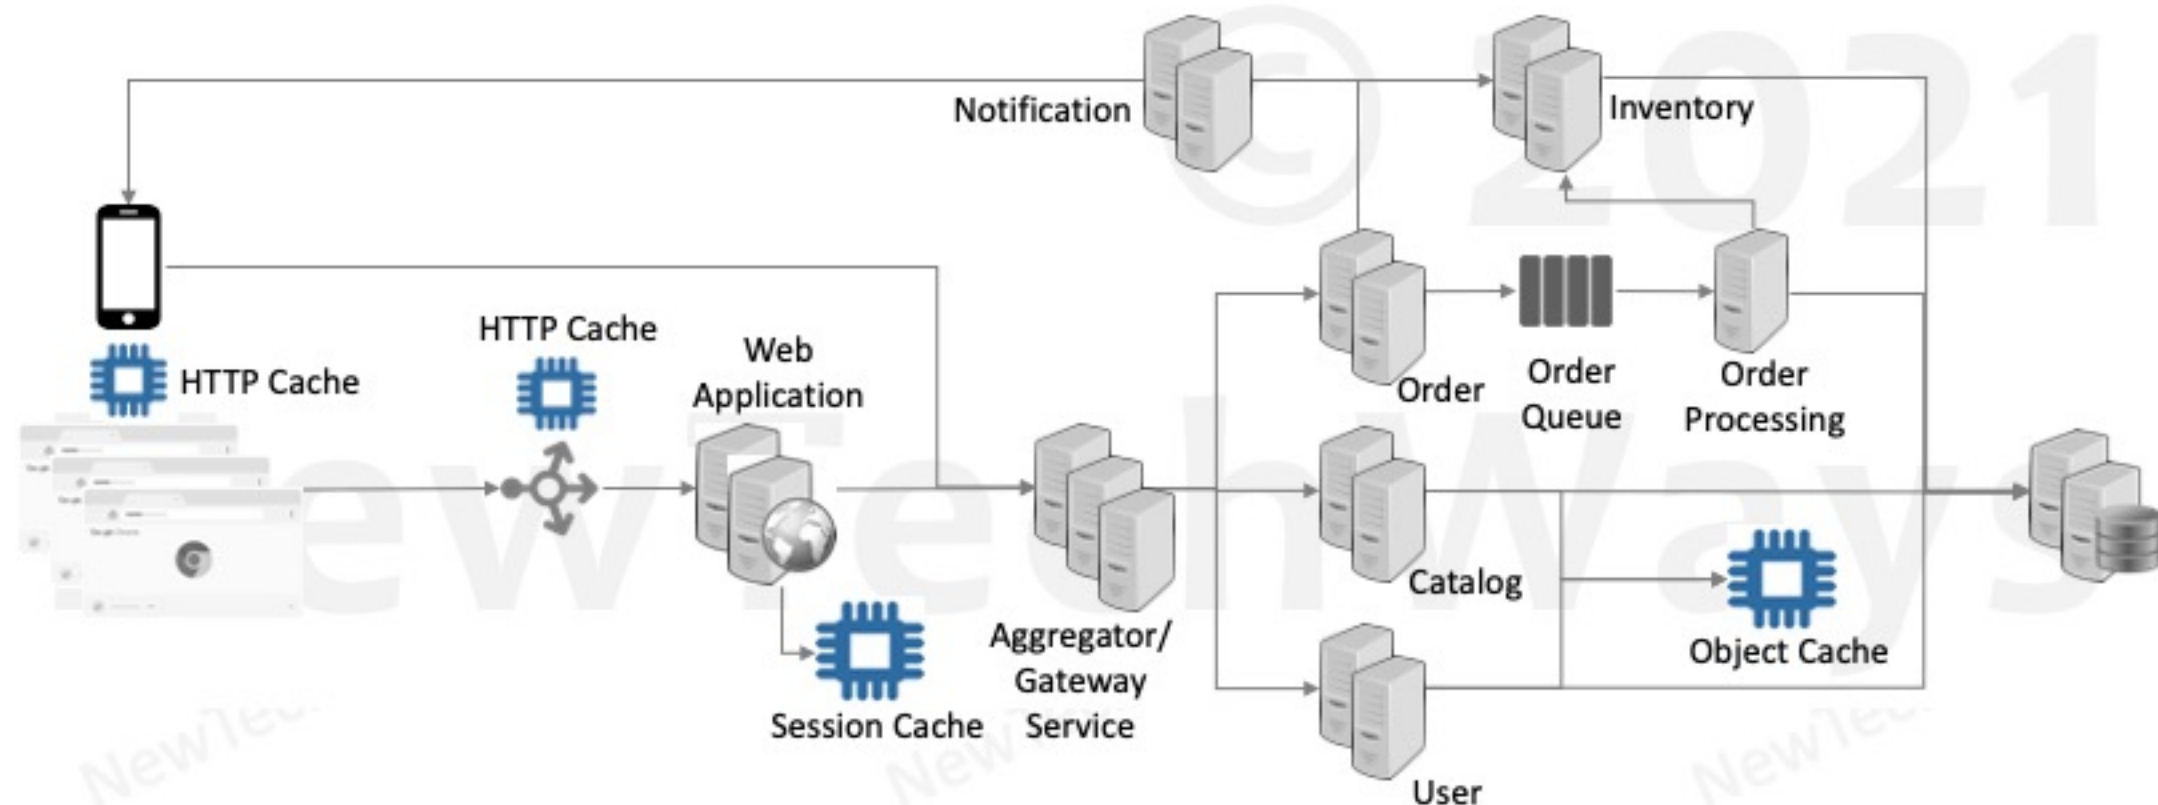
Asynchronous Processing & Scalability

- Async services require infrastructure for average load as opposed to peak load



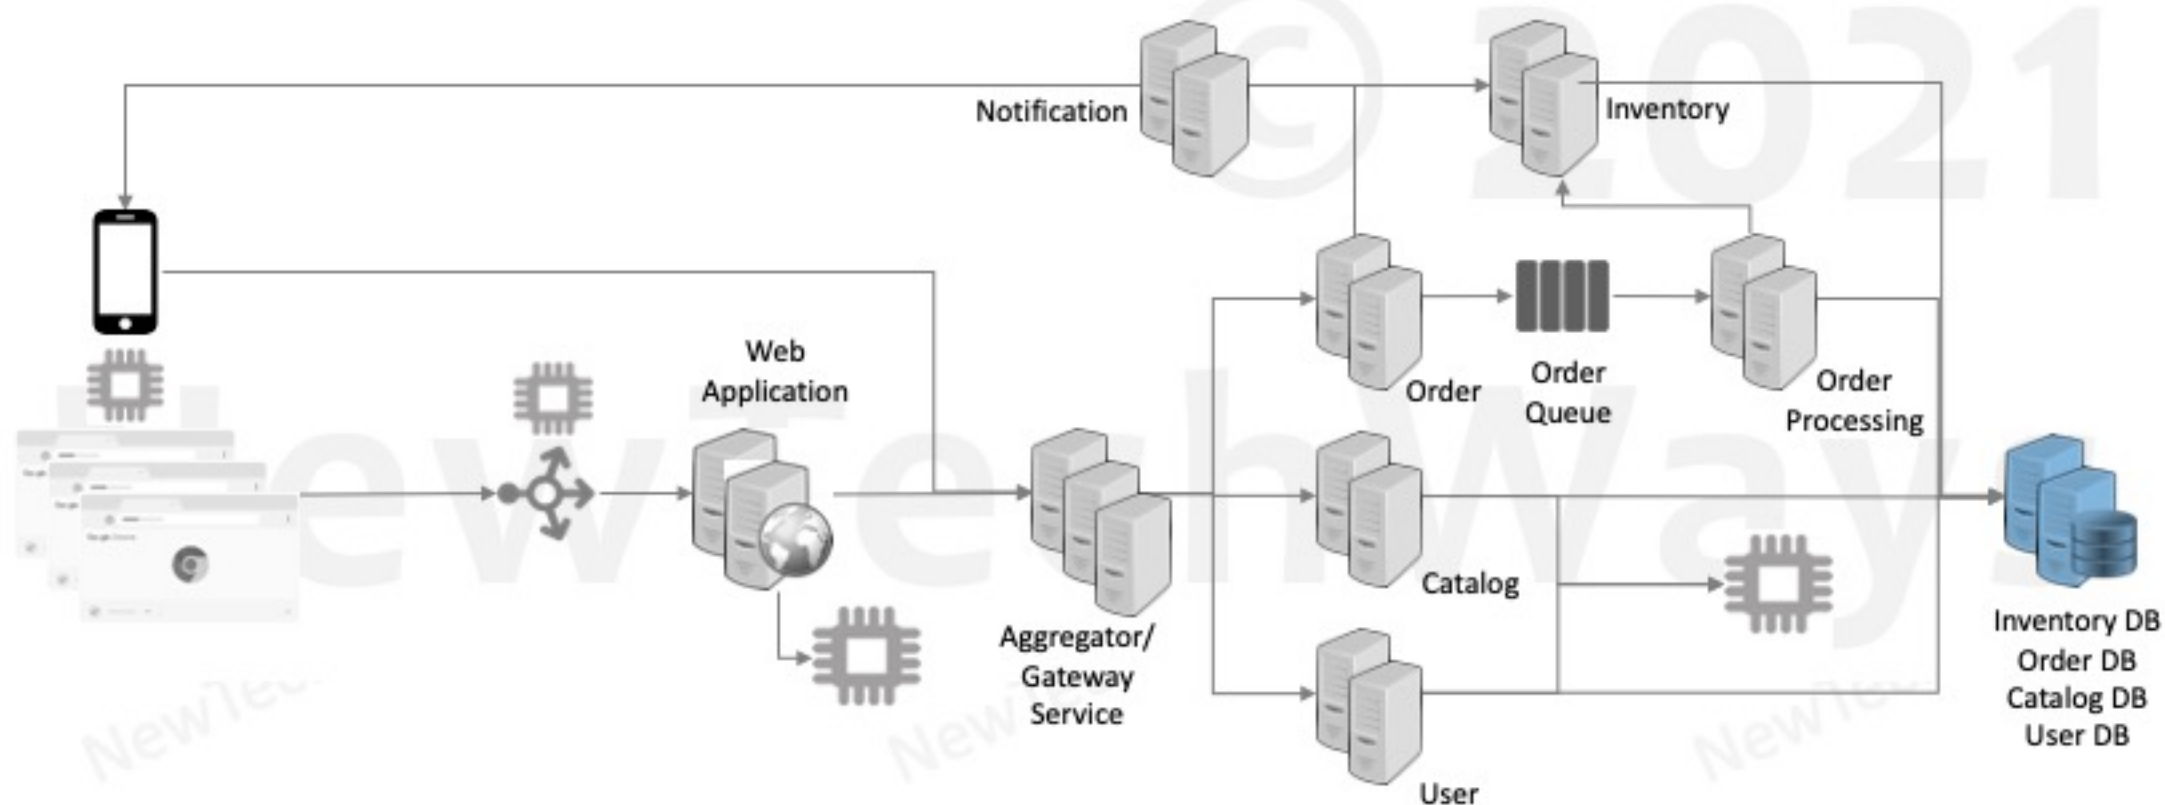
Caching

- Caching reduces latency and reduces overall read load



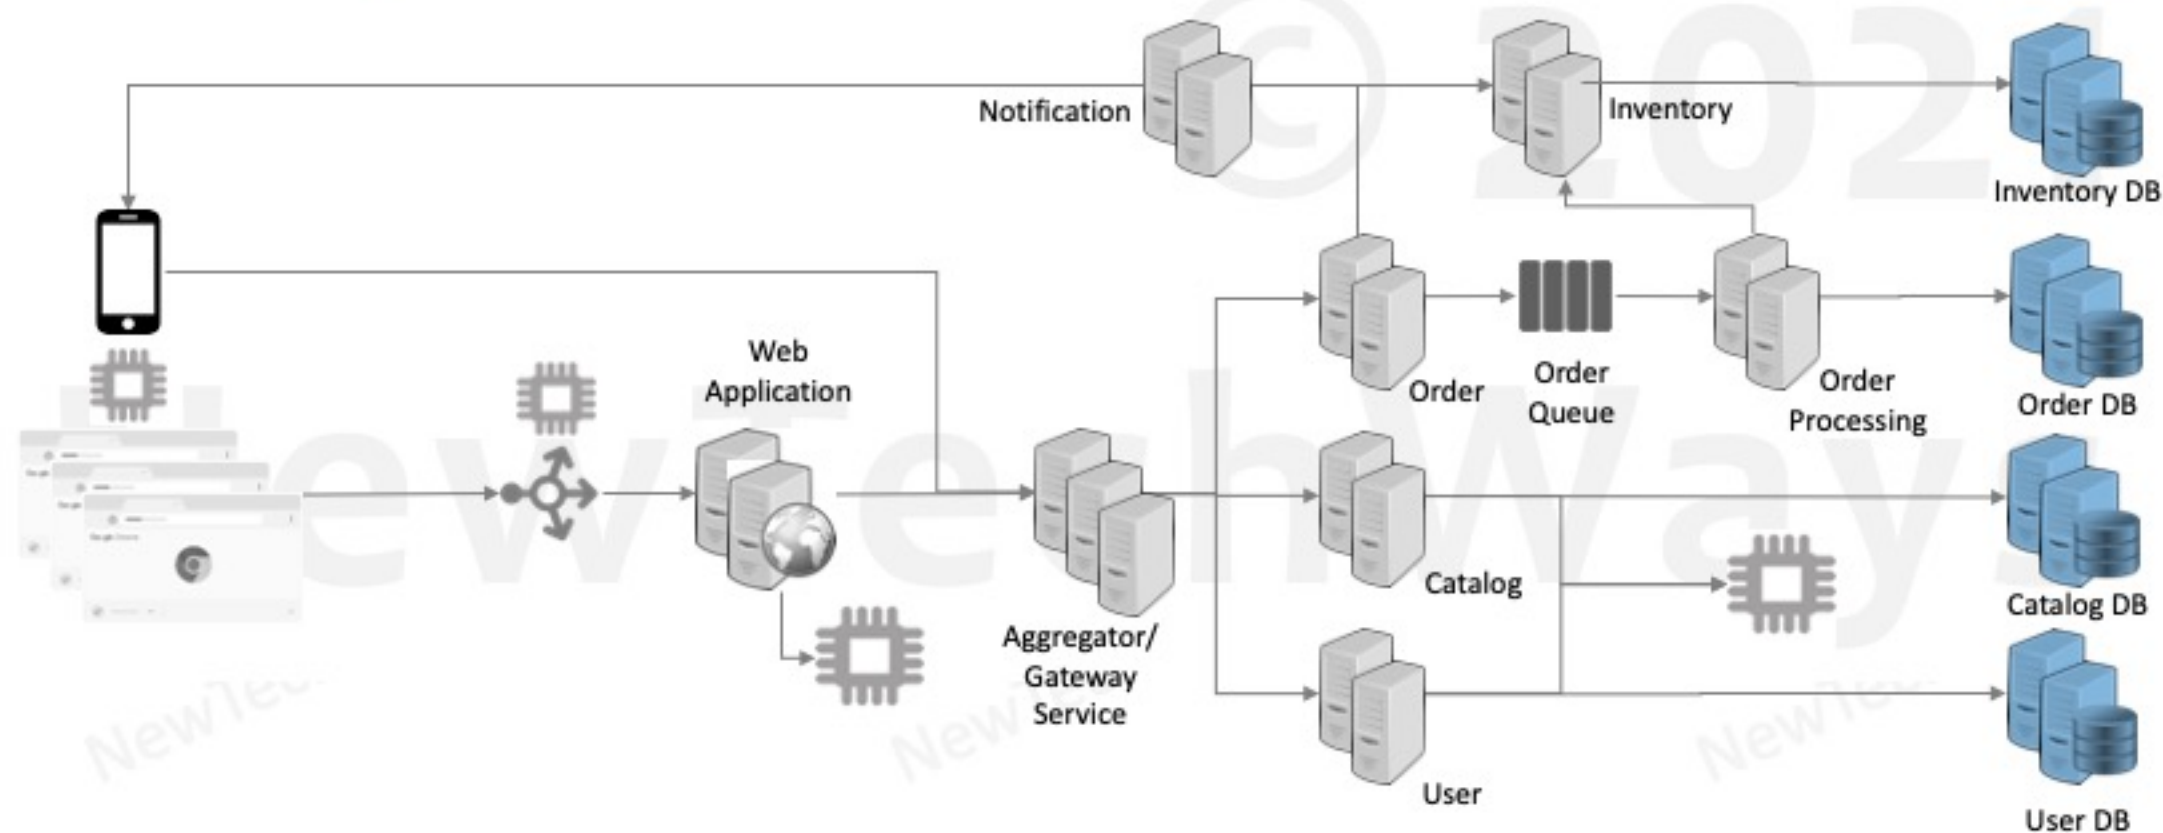
Vertical Partitioning (Micro-Services)

- Micro-Services completely decouples services and databases for higher scalability




Vertical Partitioning (Micro-Services)

- Micro-Services completely decouples services and databases for higher scalability
 - Can no longer do inter service ACID transactions and need to deal with eventual consistency

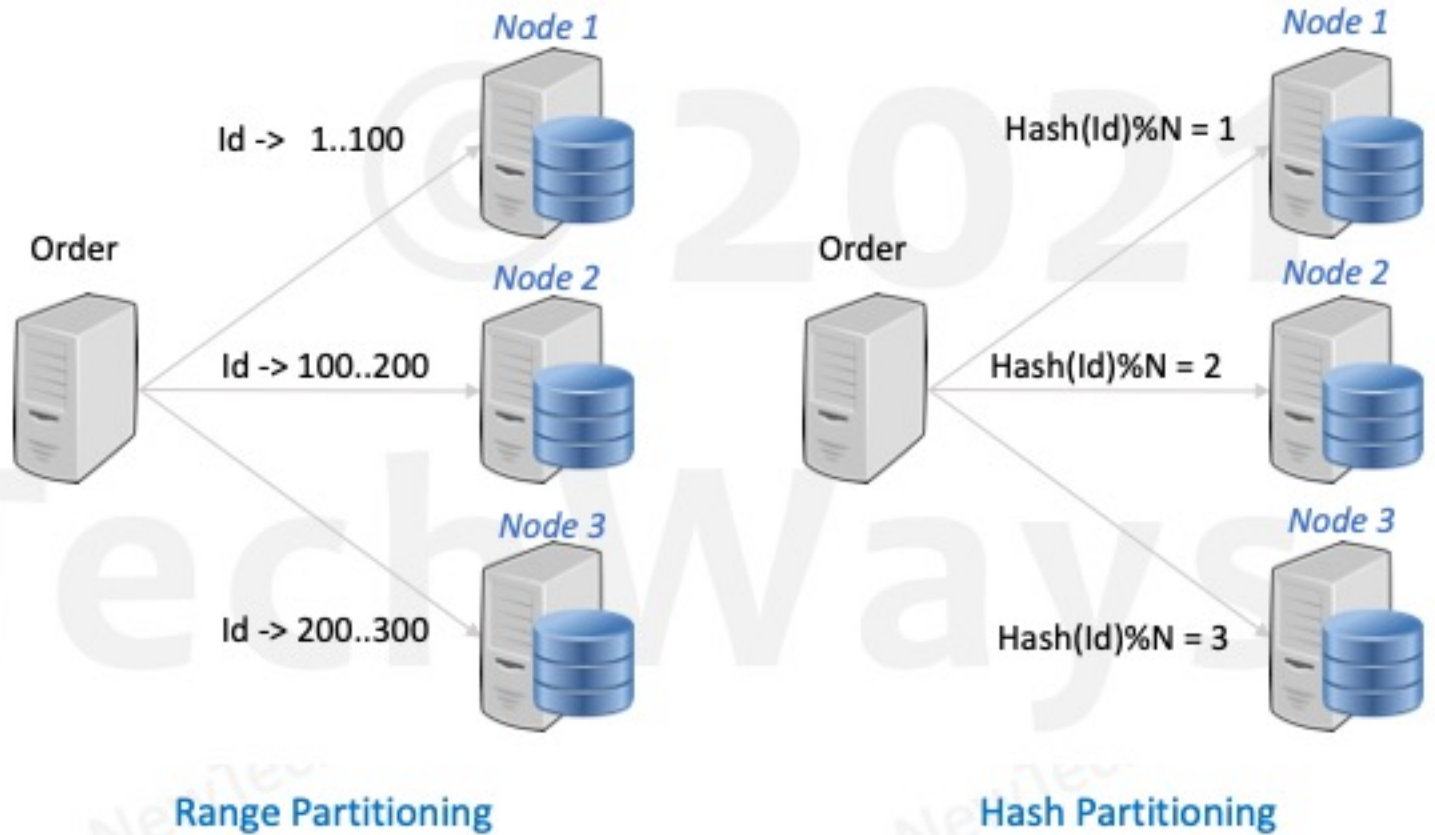


Database Partitioning

Vertical Partitioning Of System



Horizontal Partitioning Of Database




Database Partitioning Selection

Range Partitioning


`SELECT * FROM Order WHERE id = 150`

`SELECT * FROM Order WHERE id > 150 AND id < 250`




Hash Partitioning

`SELECT * FROM Order WHERE id = 150`



Routing with Database Partitioning

Get: Id = 256



Horizontal Scaling Methods

1. Services
2. Replication
 1. Stateful
 2. Stateless
3. Partitioning
 1. Vertical/Functionality Partitioning
 2. Database Partitioning
4. Asynchronous Calls
5. Caching

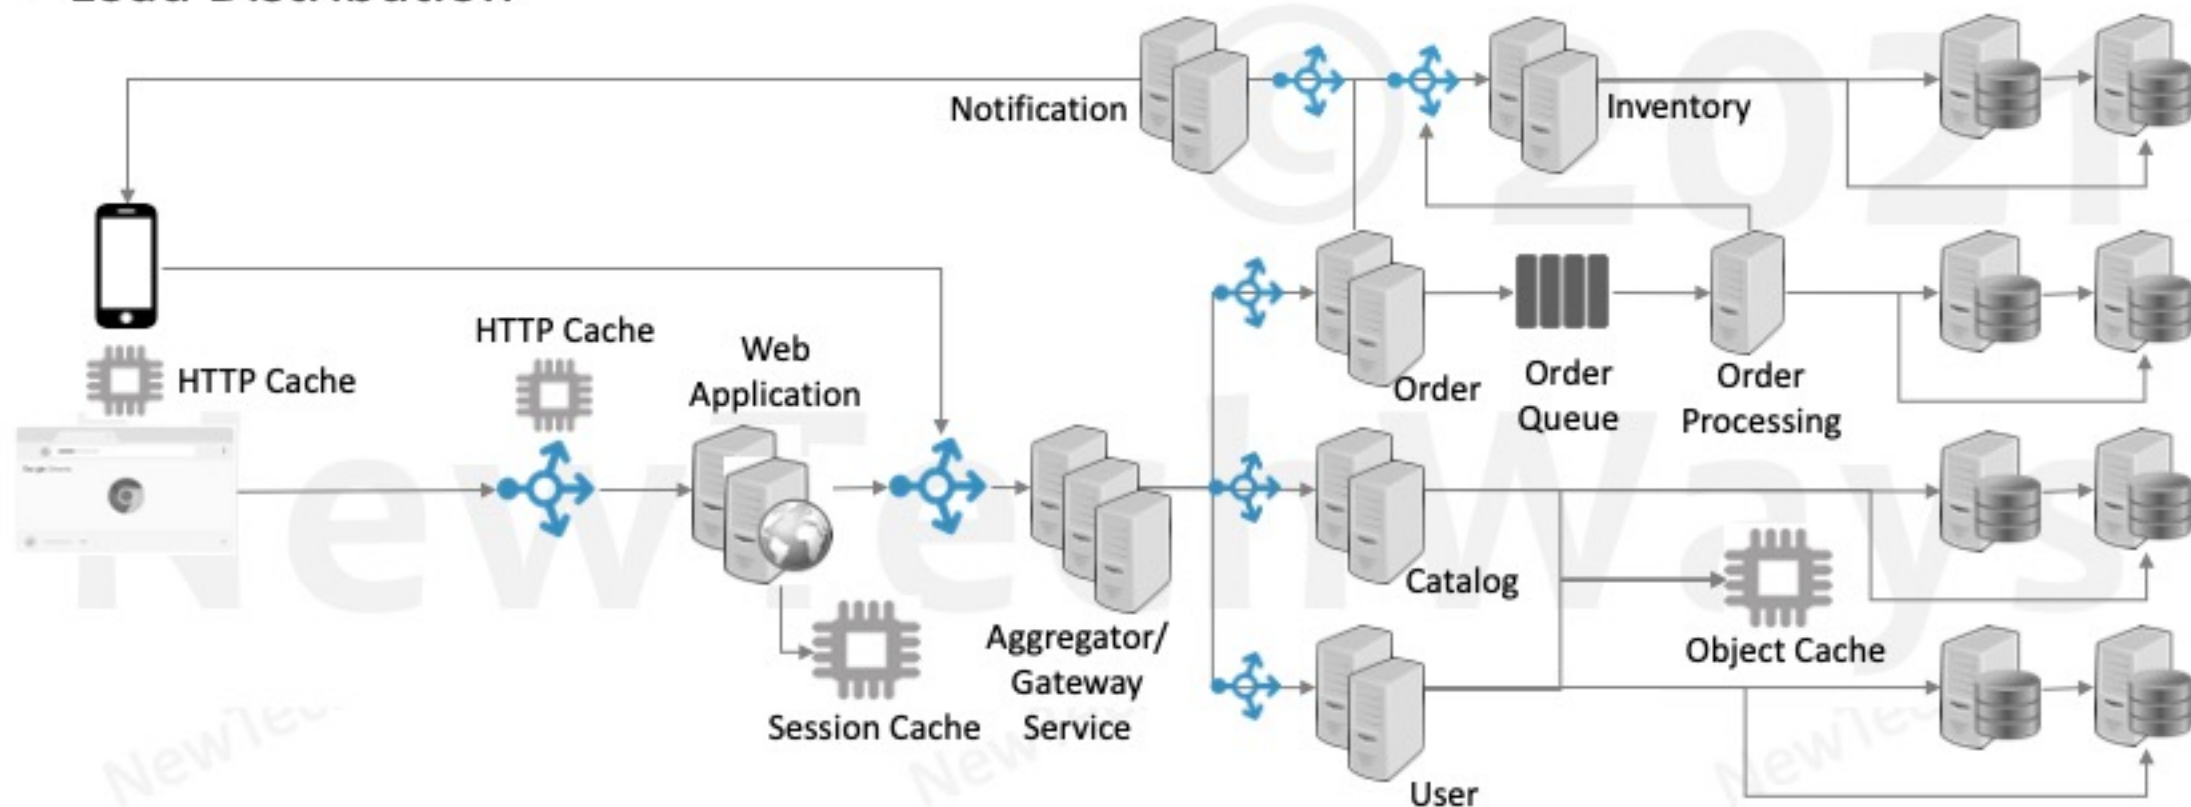
© 2021

Load Balancing

NewTechWays

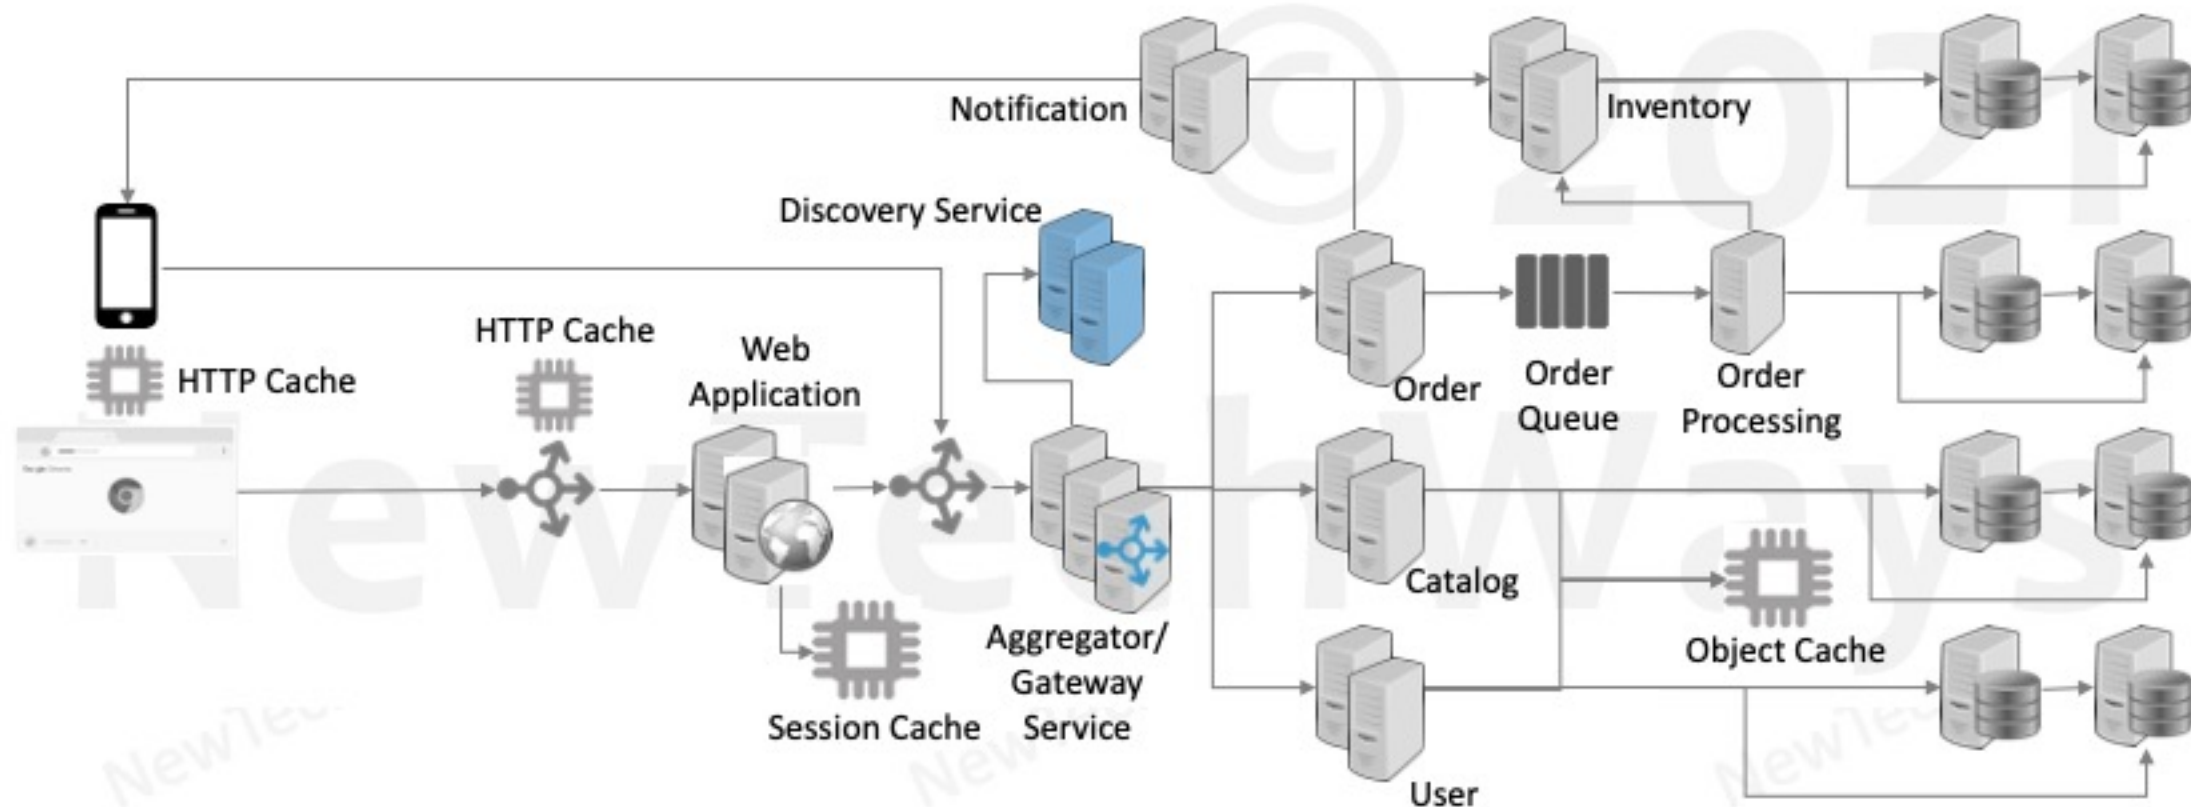
Load Balancing

- Single IP address for a Component
- Load Distribution



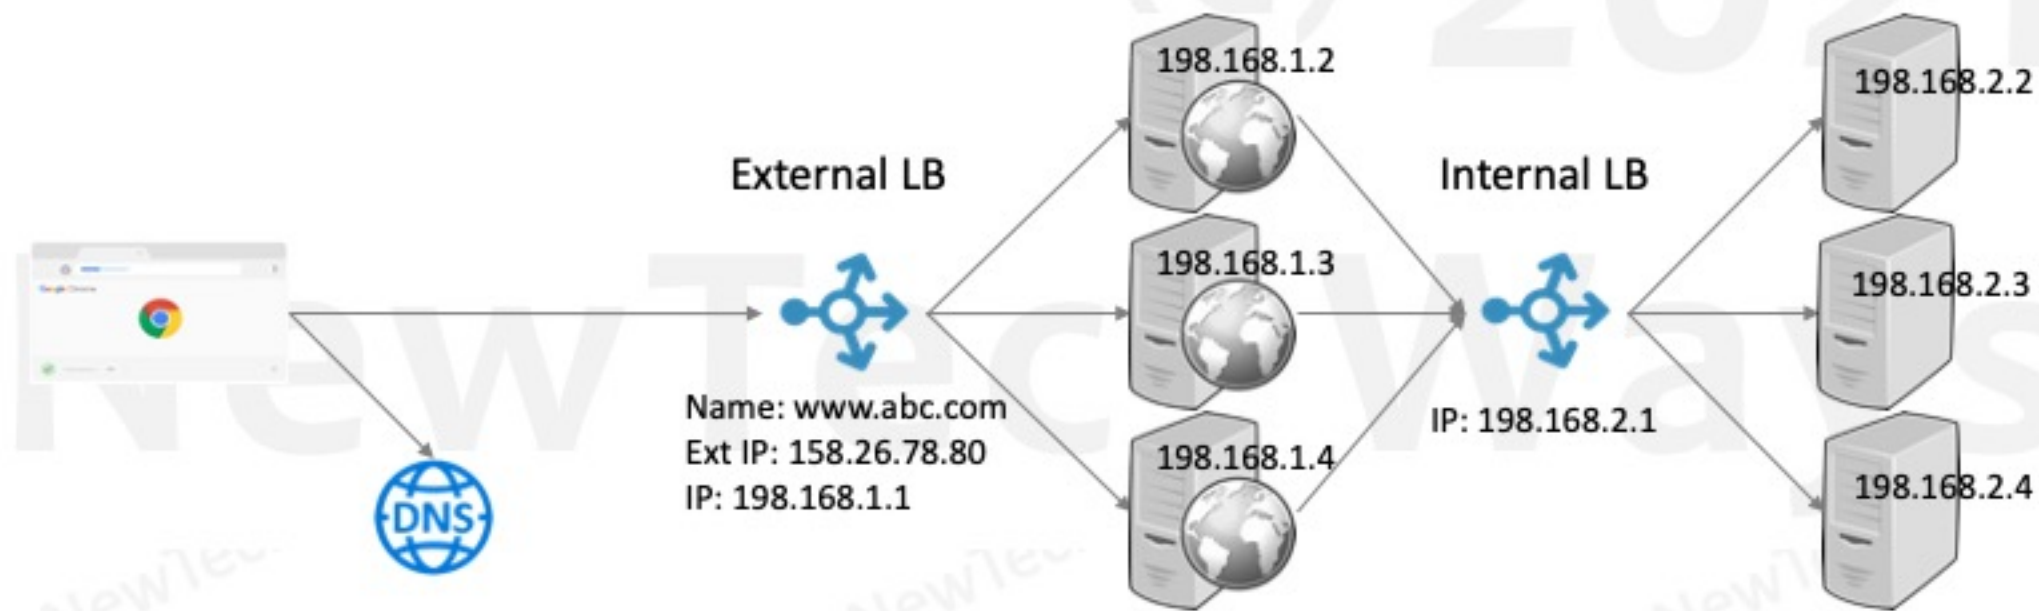
Discovery Service & Load Balancing

- Discovery – Registry for IP of Healthy Instances



Load Balancer Discovery

- External Clients – Use DNS to discover the external load balancer
- Internal Clients – Use a local registry/config to discover an internal load balancer



HLB & SLB

- Hardware Based Load Balancer
 - Load distribution for L4 & L7



2x 14-Core Intel Xeon processors
1.6TB of available use storage space
512GB DDR4 RAM
4x 100G and 8x 40G fiber ports

- F5 Big IP i5000 series
 - Connections: 300 million
 - Throughput: 320/160 Gbps
 - RPS (L7): 10 million

- Software Based Load Balancers

- Load distribution L7
- Features
 - Content based Routing
 - Supports SSL Termination
 - Supports Sticky Sessions



Apache

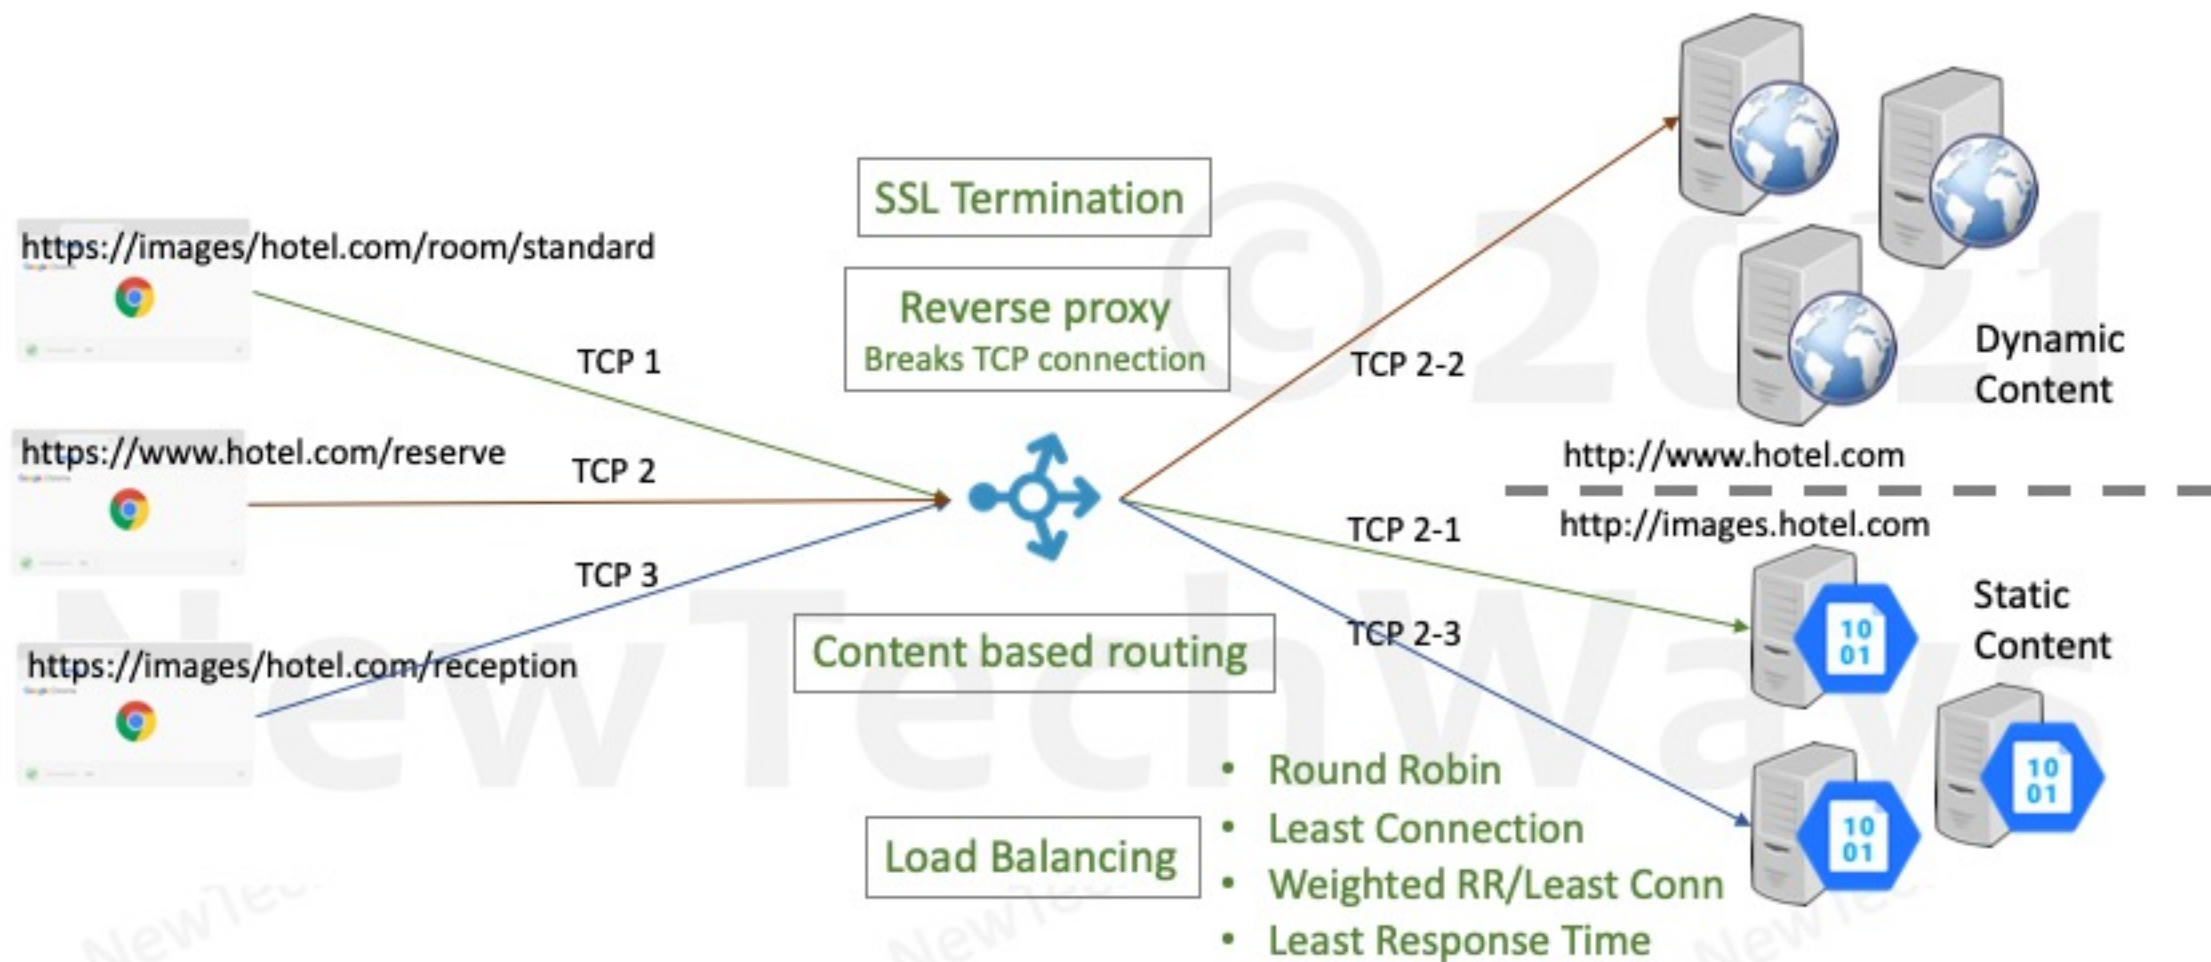
NGINX



TCP/IP Model		OSI Model	
Application Layer	HTTP, HTTPS, SMTP, IMAP, FTP, DNS, NNTP	Application	'Layer 7' Load Balancing
		Presentation	
		Session	
Transport	UDP, TCP, SCTP	Transport	'Layer 4' Load Balancing
Internet		Network	
Network Access (Link)		Data Link	
		Physical	


- NGINX
 - Connections: 225 K
 - Throughput: 70 Gbps
 - RPS: 3 million

L7 – Load Balancer



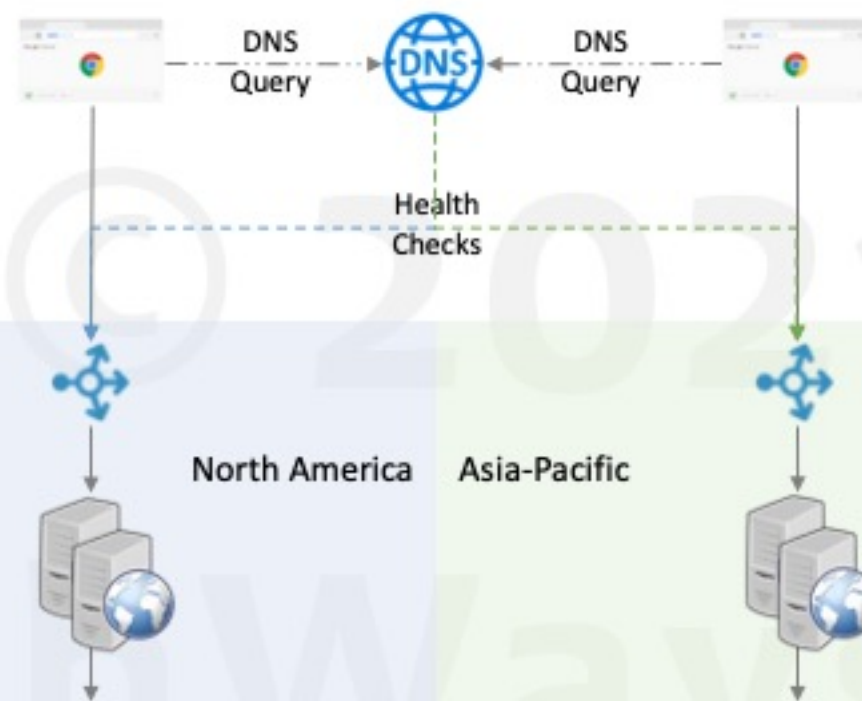
DNS as Load Balancer

- Configure DNS records with multiple A records
 - Return single IP in a round-robin fashion
 - Return a list of IP
- Cloud based DNS can be configured along with health checks
- Drawbacks
 - Indefinite caching and not respecting TTLs
 - Low or zero TTLs can create a very high load on DNS




Global Server Load Balancing

- Scalability
 - Routing for multi-geographic systems
- Performance
 - Locality for multi-geographic users
 - Client to Datacenter Latency
 - Client to Datacenter Proximity
 - Datacenter Geography
- High Availability
 - Multi region availability
- Disaster Recovery



Global Data Replication

- Active-Active Setup
 - All sites active
- Master-Master or Peer-to-Peer replication
 - Mostly asynchronous
- Failover is quick
- Some data loss is a possibility



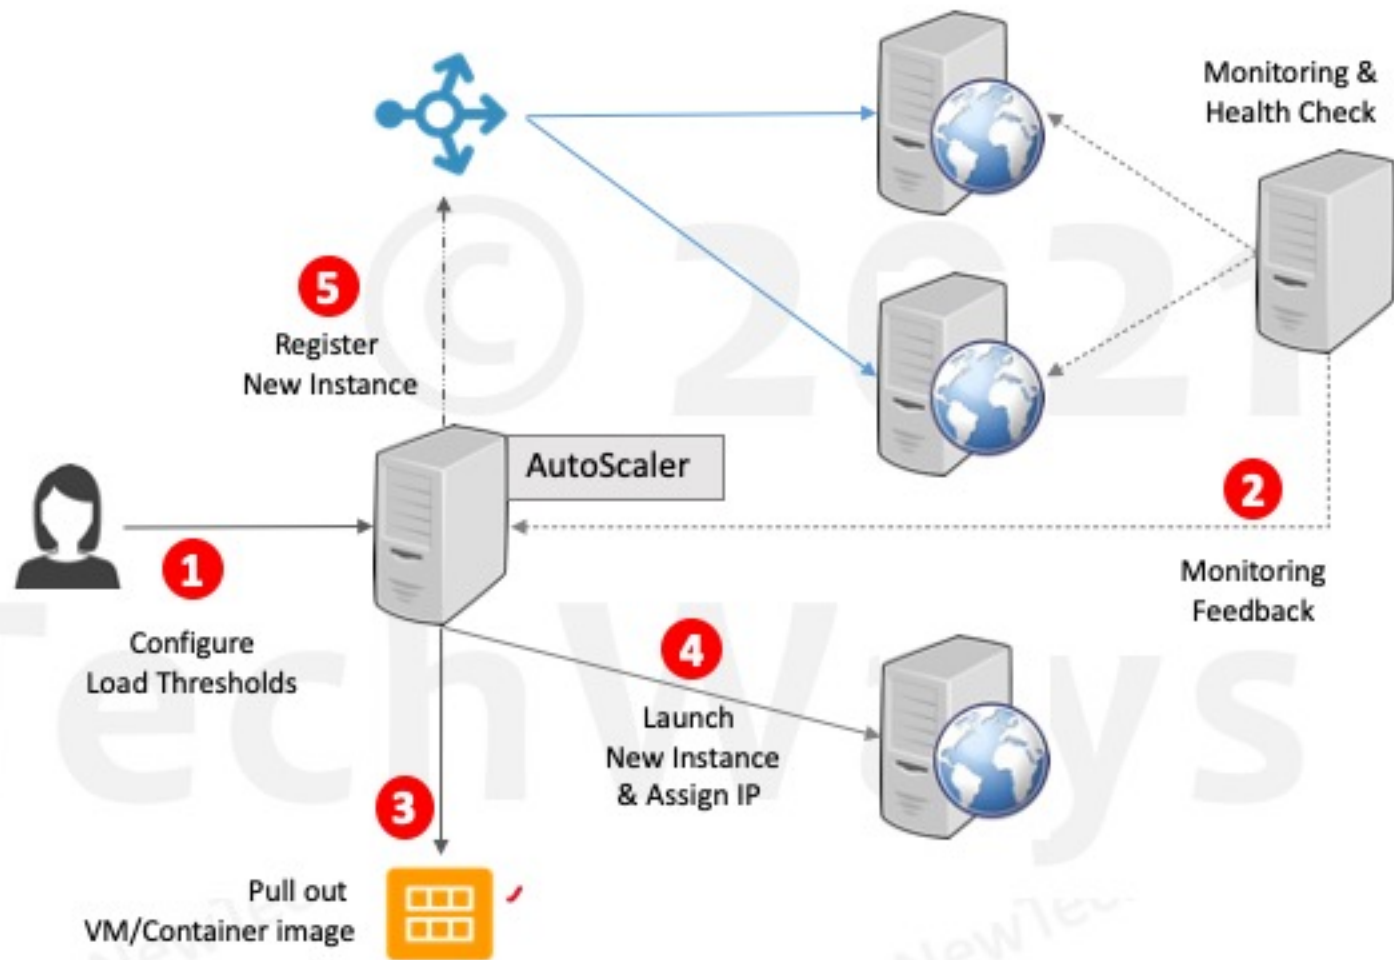
Auto Scaling

- **Monitoring Service**

- Monitor Load
 - CPU
 - Network
 - Disk
- Monitor Health
 - Ping
 - Http

- **Auto-Scaling Service**

- Configure load thresholds
- Monitor load
- Launch New Instance
- Shutdown Instance



© 2021


Micro-Services Architecture

NewTechWays

Micro-Services Motivation

High Scalability

Decentralization & Independence



Frequent Deployment


Independent Deployment

Independent Development

Independent Services

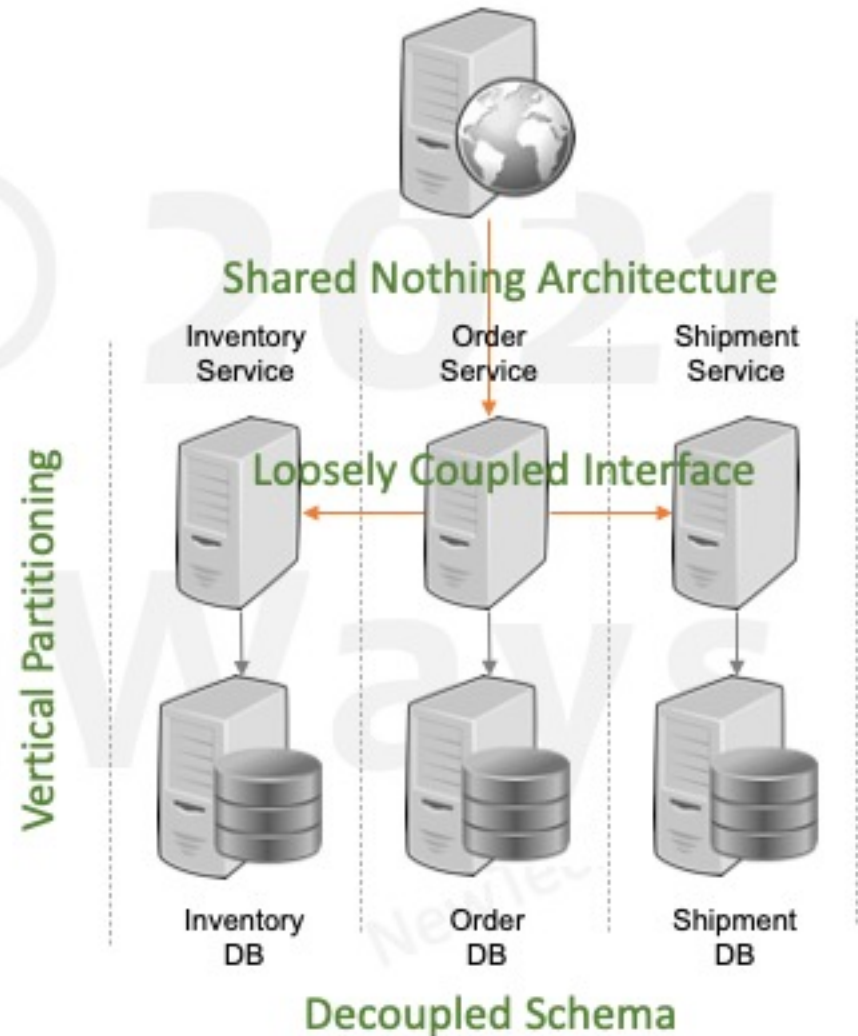
Service Oriented Architecture

- Independent
 - Each service can have its own technology stack, libraries, frameworks etc.
 - Each service can be scaled independently and differently
- Not Independent
 - Common interface schema
 - XML schema
 - Common database schema
 - RDBMS schema
- Issues
 - Service development may be independent but not deployment
 - Single database has scalability limitations



Micro-Services Architecture Style

- Shared Nothing Architecture
 - Services developed and deployed independently
 - Achieved through vertical partitioning
- Vertical/Domain Partitioning
 - Independent schema/database
 - Loosely coupled service interfaces
 - REST interfaces instead of XML/WSDL schemas
 - No reusable libraries except utilities
- Issues
 - Duplicate Codebase
 - Transaction failures
 - Transaction rollbacks



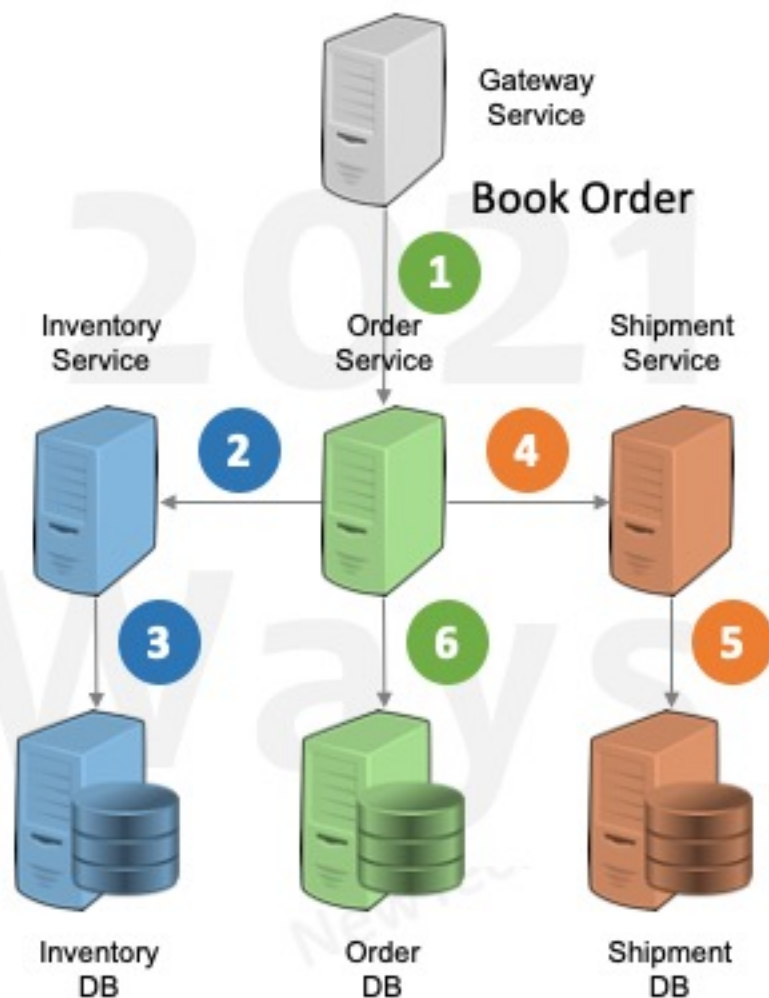
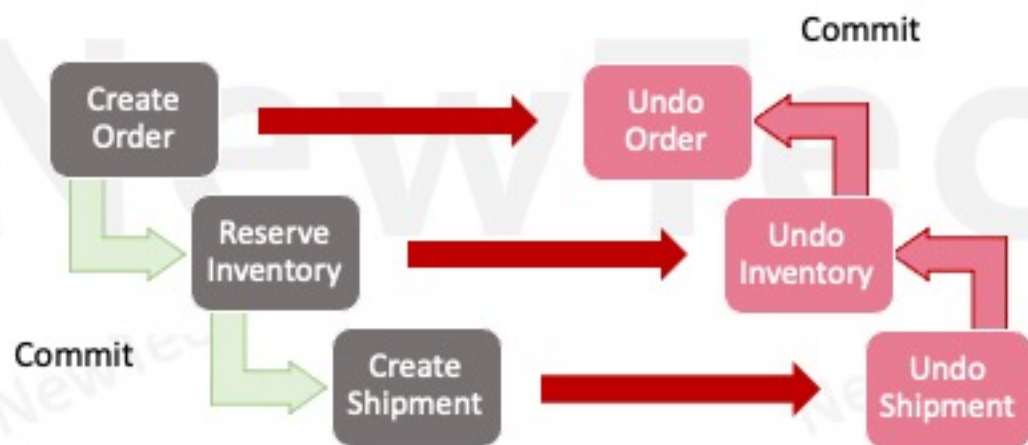
Micro-Services Transactions

- Transaction involves multiple machines
 - Distributed services with their own DB
 - Local transaction not possible
- Options
 - Distributed ACID Transactions
 - 2PC/3PC
 - Completely ACID
 - Compensating Transactions
 - SAGA pattern
 - Eventually consistent model
 - Relaxes consistency and isolation




Compensating Transactions – SAGA Pattern

- 'Logical Undo' of a partially committed transactions
 - Flow of reversal may not be exactly opposite, and some steps can be executed in parallel
 - Compensation itself can fail. Should be able to restart itself and retry
- Asynchronous processing for reliability

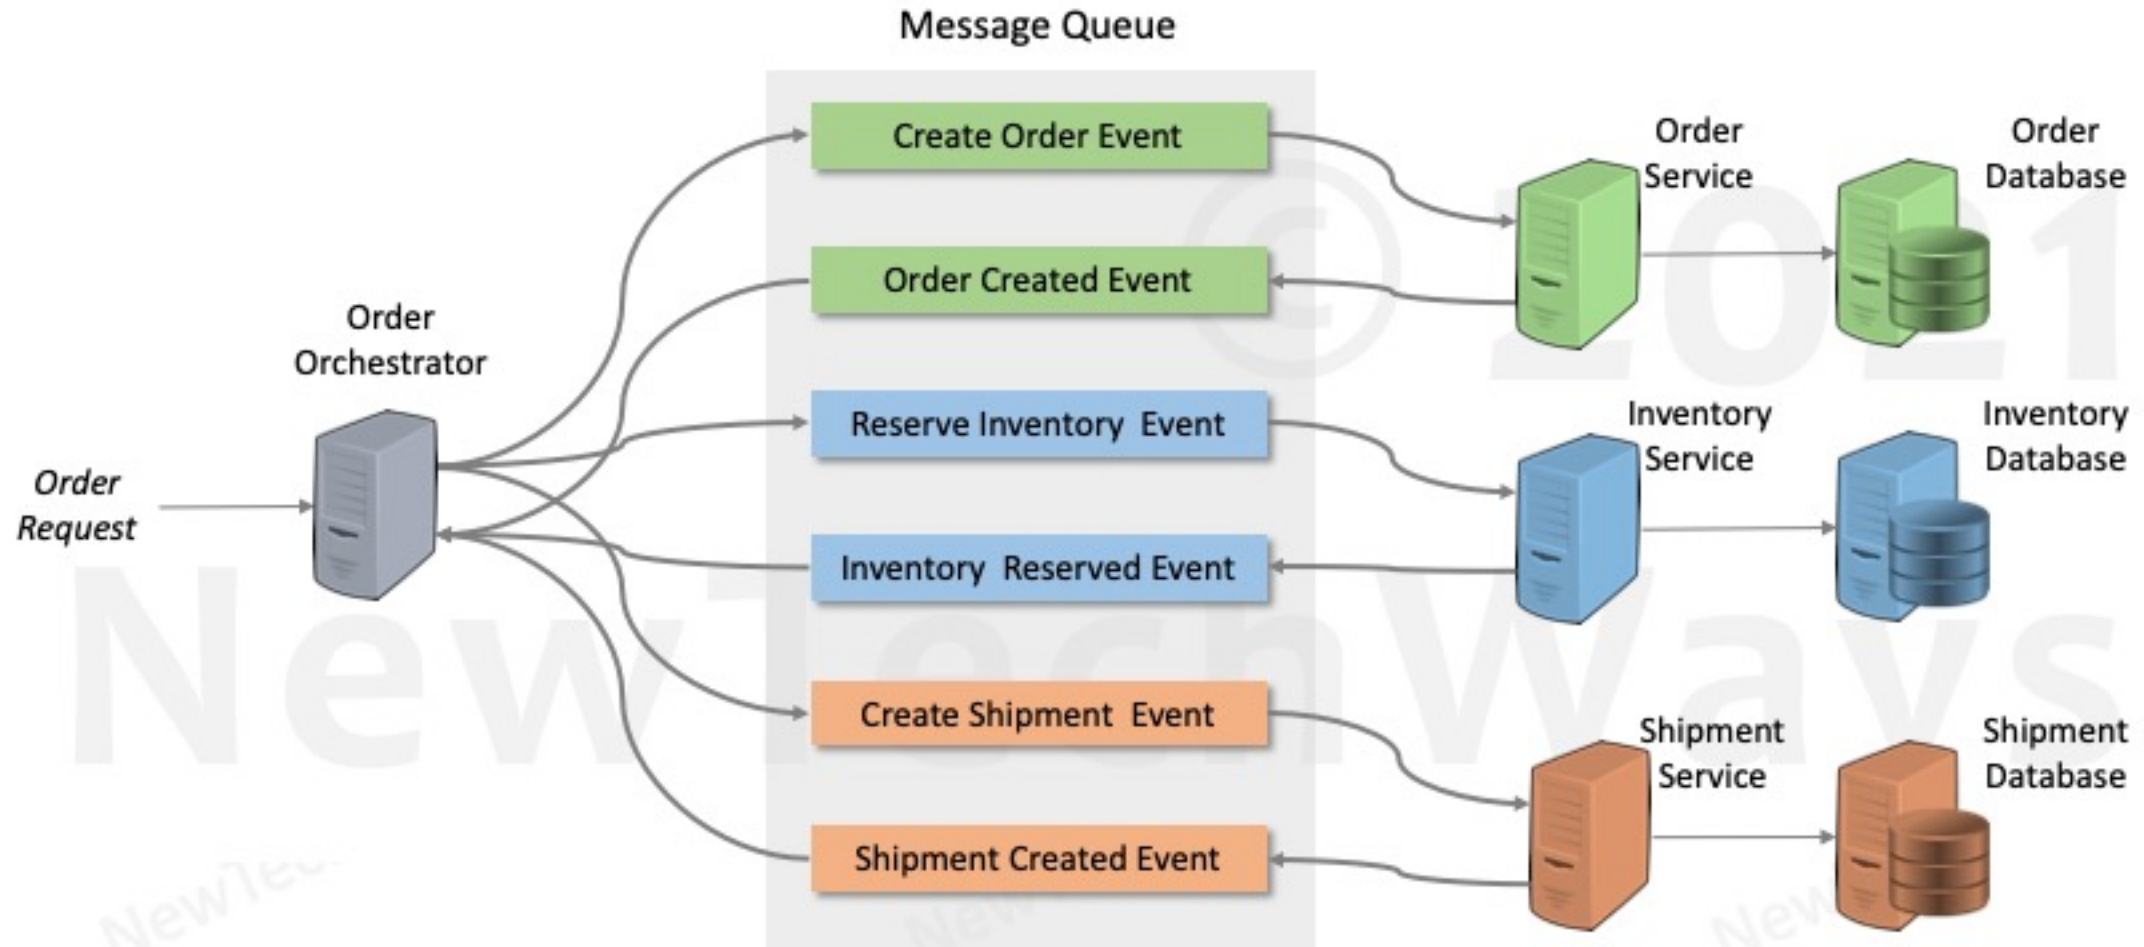


Micro-Services Communication Model

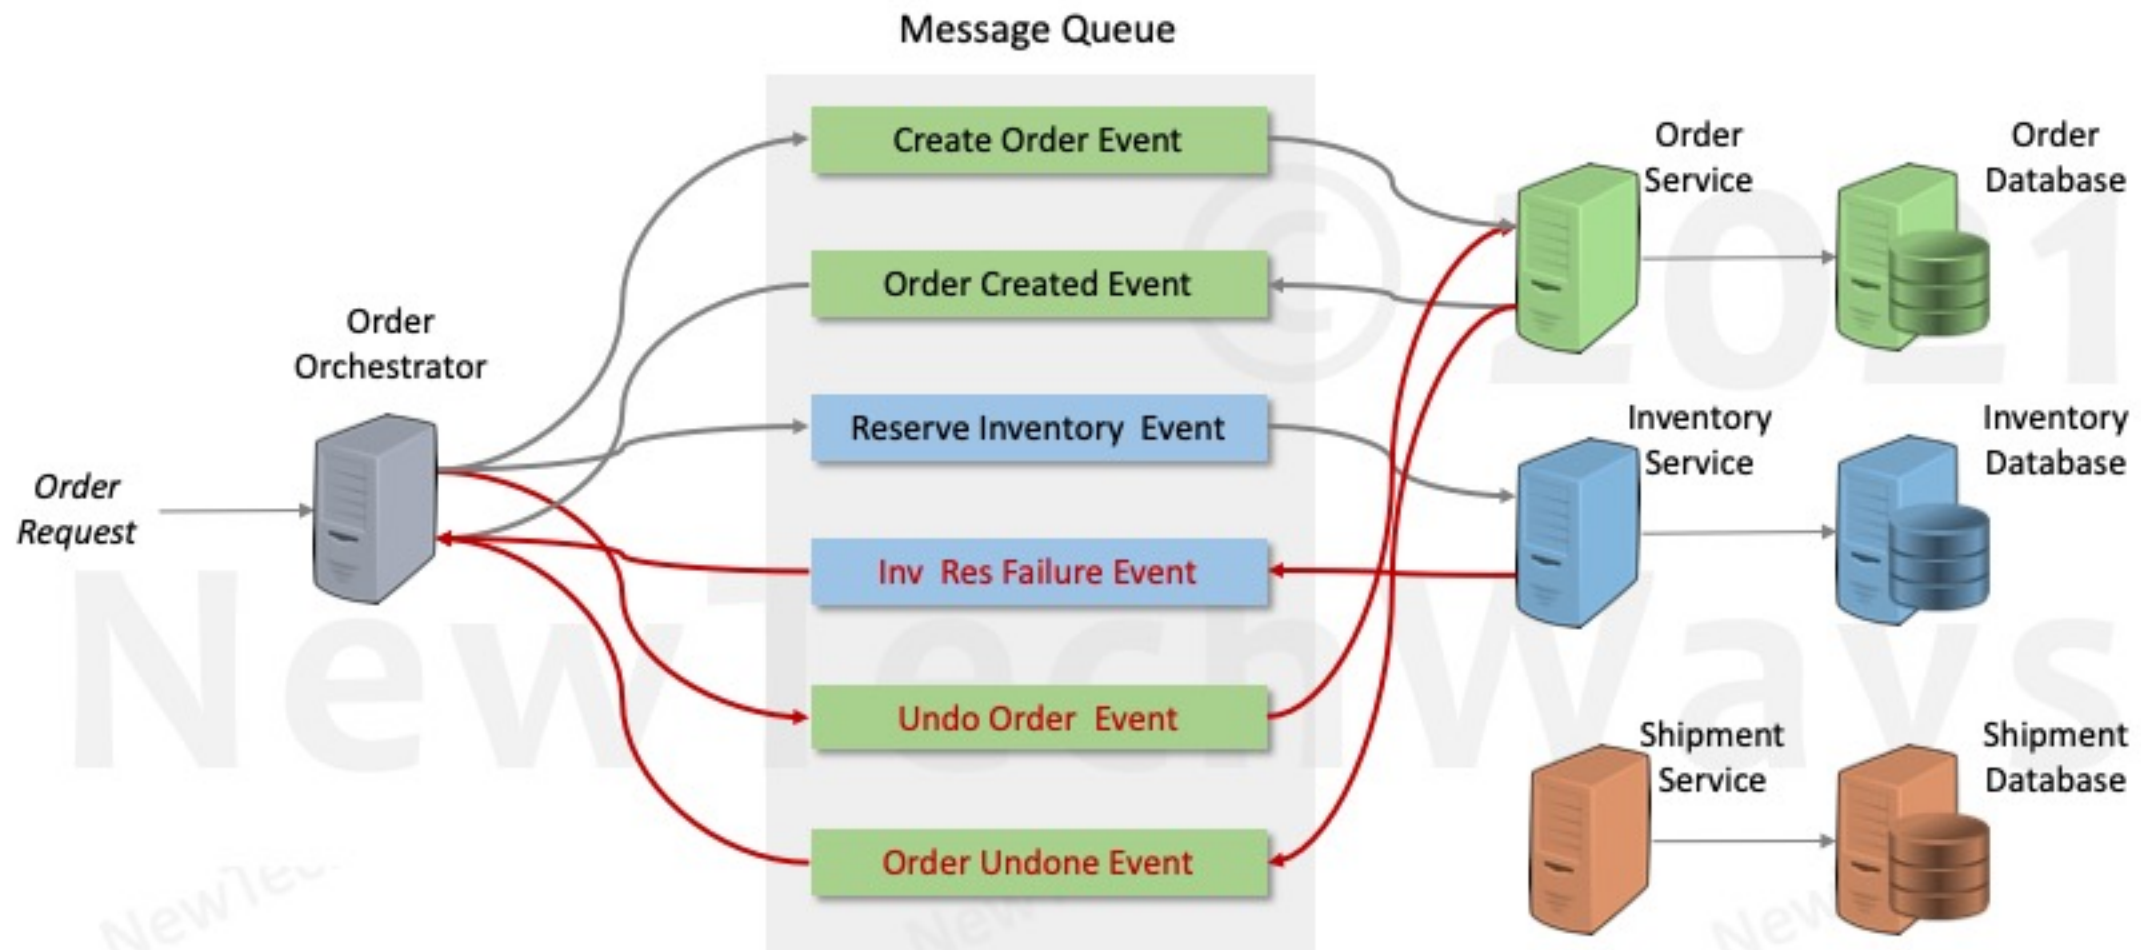
- Synchronous processing
 - Immediate Response
 - For read/query loads
- Asynchronous processing
 - Deferred response
 - For write/transaction loads
 - Higher scalability
 - Higher reliability



Micro-Services Event Driven Transactions

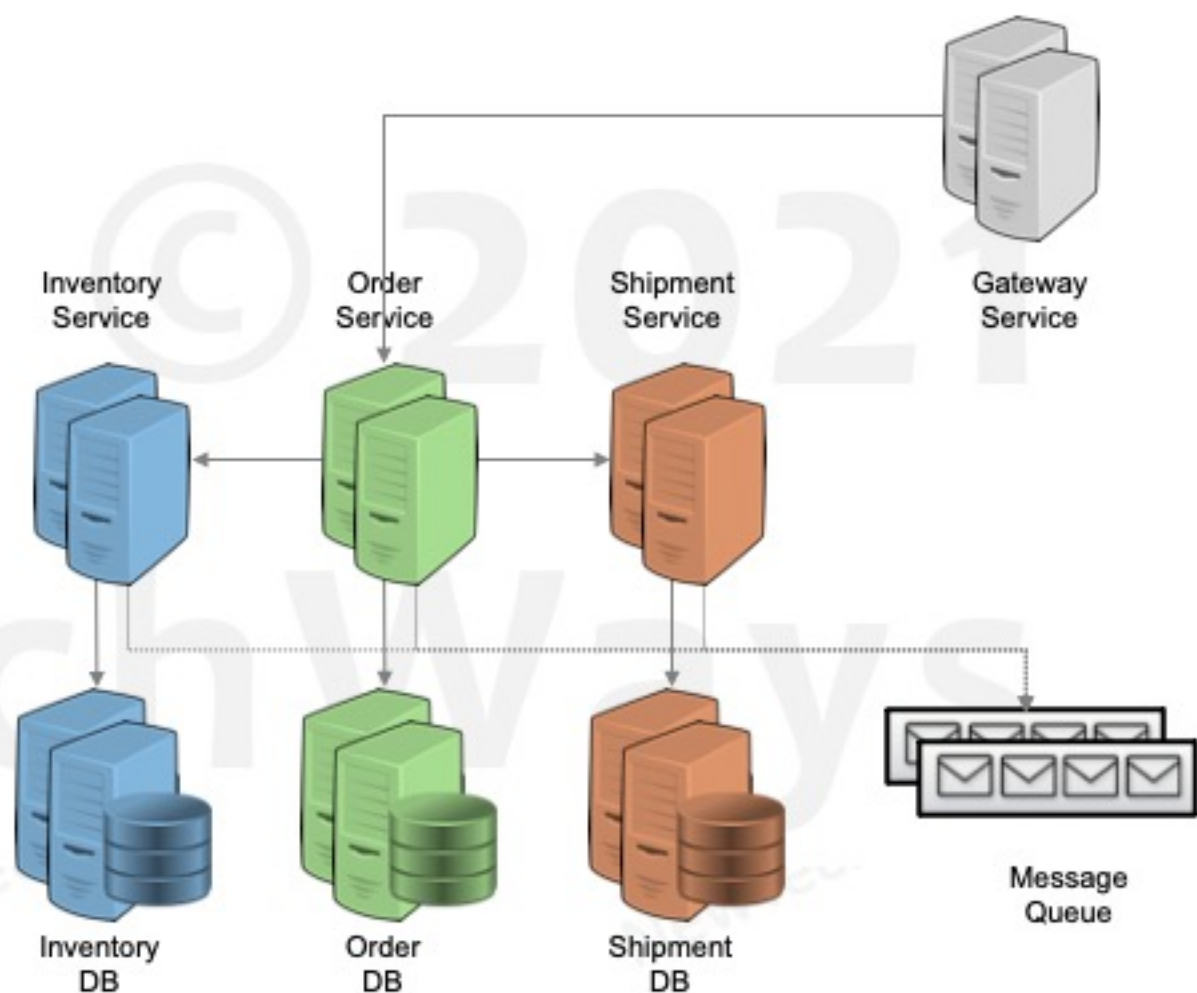


Micro-Services Event Driven Transactions



Extreme Scalability With NoSQL & Kafka

- **Microservices Transactions**
 - ACID within Service
 - Compensating Transaction across services
 - Eventually consistent
- **NoSQL DB**
 - ACID transaction at aggregate level
 - Eventually consistent transactions across aggregates
 - Low latency operations
 - Multiple nodes
 - High scalability
 - Horizontal partitioning
- **Kafka**
 - Horizontal partitioning of topics



Summary

- Scalable systems are decentralized, and their components function independently
- To make a system scalable:
 - Cache frequently read and rarely mutating data to reduce load on the backend
 - Asynchronous or Event driven processing for distributing load over time
 - Vertical partitioning of functionality into independent, stateless, replicated services
 - Partitioning and replication of state for extreme scalability
- Scalable systems requires infrastructure:
 - Load balancers – Hardware based (L4 + L7) & Software based (L7)
 - Discovery services for service discovery and health checks
 - DNS as load balancers at global scale
- Microservices for extreme scalability
 - Fully vertically partitioned services and databases leads to eventual consistency

Thanks!



NewTechWays

<https://www.newtechways.com>